



ACADGILD

Session 3: HDFS Internals

Assignment 2

Assignment 2 – Try the given quiz questions and provide the answers in a word document.

Table of Contents

1. Introduction	3
2. Objective	3
3. Prerequisites:	3
4. Associated Data Files	3
5. Problem Statement:	3
7. Approximate Time to Complete Task	5

1. Introduction

In this assignment you need to select one right choice for the questions given on the topics discussed.

2. Objective

This assignment will help you to consolidate the concepts learnt in the session 3.

3. Prerequisites:

None

4. Associated Data Files

None

5. Problem Statement:

1. Which is not the property of hdfs-site.xml file?

- a. Block size
- b. Replication factor
- c. Secondary NameNode port address
- d. Application server

Answer : b

2. What are the properties that we can edit in hdfs-site.xml file?

- a. Block size
- b. Replication factor
- c. Block reporting interval
- d. All the above

Answer : d

3. Underlying storage layers where MapReduce programs are written:

- a. Abstracted
- b. Open to all
- c. Locked
- d. None of these

Answer : a

3. Which is not a valid method in FSDATA Output Stream?

- a. close()
- b. open()
- c. getpos()
- d. sync()

Answer : a

5. How many blocks of size 128MB will be allocated for a file of size 524288KB:

- a. 2
- b. 3
- c. 4
- d. 5

Answer : c

6. MapReduce Job client calculates the input split by:

- a. Figuring the first and last whole records in the block
- b. Figuring only the first block
- c. Figuring only the last block
- d. None of these

Answer : a

7. Mappers are directly related to:

- a. Input data
- b. Input splits
- c. Output data
- d. Output splits

Answer : b

8. After data is written in an HDFS file, HDFS does not provide any guarantee that data are visible to a new reader until the file is _____

- a. closed
- b. opened
- c. writing

Answer : a

9. What mechanism does Hadoop use to make namenode resilient to failure?

- a. Take backup of filesystem metadata to a local disk and a remote NFS mount
- b. Store the filesystem metadata in cloud
- c. Use a machine with at least 12 CPUs
- d. Using expensive and reliable hardware

10. All the files in a directory in HDFS can be merged together using:

- a. getmerge
- b. putmerge
- c. remerge
- d. mergeall

Answer : d

6. Expected Output

None

7. Approximate Time to Complete Task

15 mins